Introduction to parallel computing

Jeremy Iverson

College of Saint Benedict & Saint John's University

cpu performance

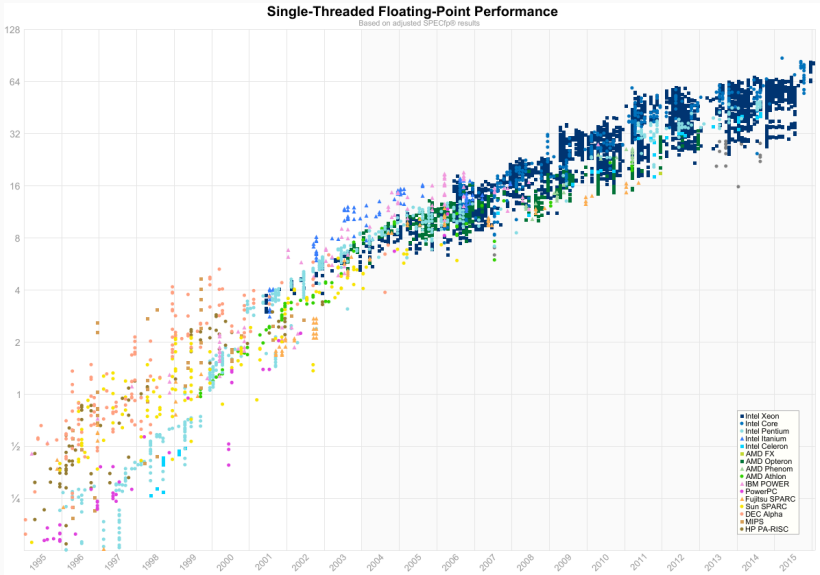

Single-Threaded Floating-point Performance by HenkPoley

- the “almost” current state of things
- what do you notice about this graph?
  - from 1986 – 2002, microprocessor performance increased by roughly 50% every two years, since then it has dropped to about 20%

clock speed

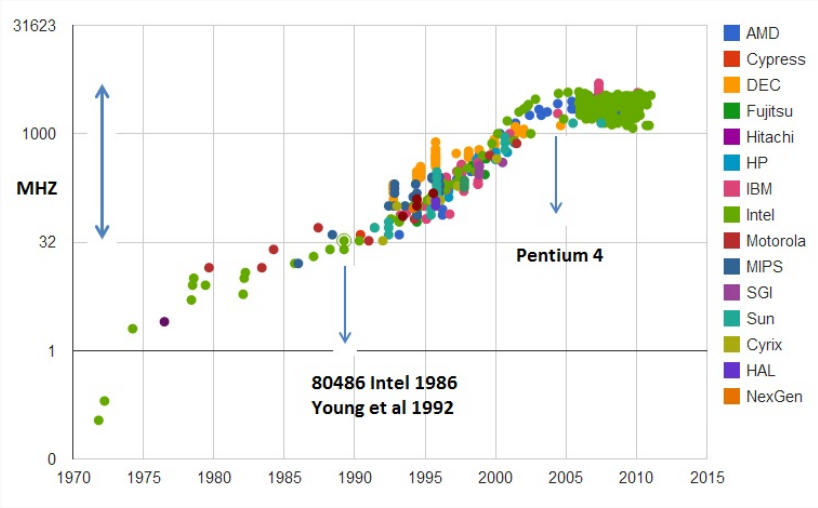

Clock CPU Scaling.jpg by Newhorizons msk / CC0 1.0

- do you notice anything odd about this graph, given the previous graph?
  - we are seeing ~20% performance improvement despite stagnating clock speed
- why does clock speed stop increasing?
  - faster processor = increased power consumption.
  - increased power consumption = increased heat.
  - increased heat = unreliable processors.
- we are reaching / have reached the limit that current cooling technologies can remove heat from the chip (a.k.a. the power wall)
- so how does performance improve, when clock speed plateaus?

transistor count

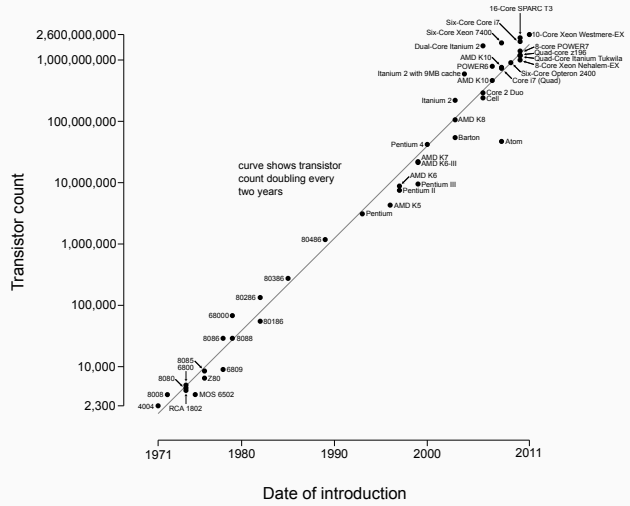

Transistor Count and Moore's Law - 2011 by Wgsimon / CC BY 3.0 / cropped from original

- pre 2005 — a golden era when clock speed could be increased while adding more transistors, which resulted in the large performance improvements in the first graph
- post 2005 — performance improvements now have to come from more advanced chip capabilities
- we don't speed up the processor, we make it more sophisticated
  - larger data busses
  - larger caches
  - deeper instruction pipelines (ILP)
- Moore's law (c. 1965) — transistor count roughly doubles every two years
- so, how can we keep continue to improve performance?

supercomputer performance

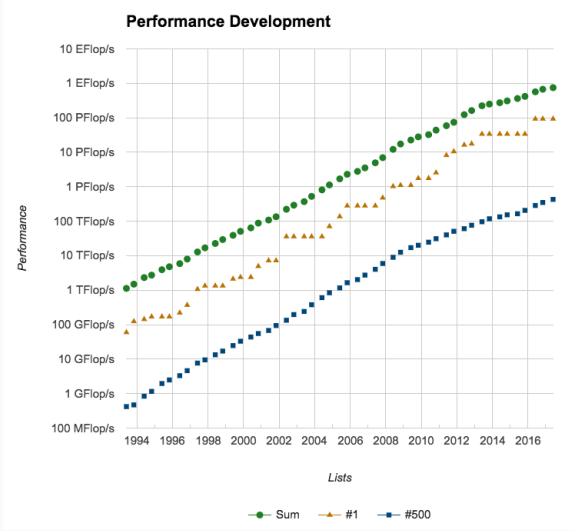

Graph courtesy of The Top 500 List.

- so how does any of this affect us, in this class?
- notice anything about this graph?
  - performance improvement is exponential
  - how can that be?
    - we are no longer constrained by the limitations of a single chip
    - so what are our limitations?

what are the challenges of parallel computing?

- how do we take advantage of available parallel resources, i.e., how to convert serial programs to parallel programs
  - automatic translation has been attempted and is still being studied, but success has been limited

counting

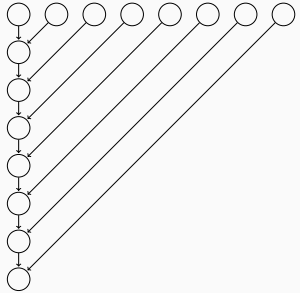

- how many steps, described as a function of the number of nodes?
- where are there dependencies / why can't this be parallelized?

```
for (int i=0; i<n; ++i) {
  a[0] += a[i];
}
```

counting cont.

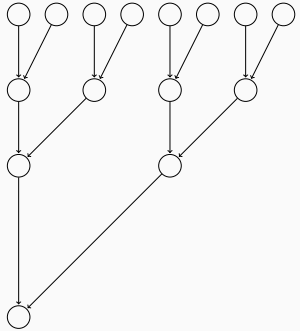

- how many steps, described as a function of the number of nodes?
- what happens if the number of nodes is not a power of two?
- what would the code look like to accomplish this (what do you need to express, that you do not know how)?

except where otherwise noted, this worked is licensed under creative commons attribution-sharealike 4.0 international license