

Computer organization

Jeremy Iverson

College of Saint Benedict & Saint John's University

plan for the day


- computer organization
 - central processing unit
 - main memory
 - cache memory
- case study: matrix multiplication
- using Git
- `printf("hello, world\n");`

central processing unit (CPU)

what does memory look like?

have students draw a picture of “memory”

memory organization



- memory is organized sequentially

- a process is an instance of a computer program
- a multitasking operating system is capable of running many processes despite having only one or a “small” number of physical cores
 - context switches the processes at predefined time slices
 - **top** activity


memory organization

0x0000	
0x0001	
0x0002	
0x0003	
0x0004	
0x0005	
0x0006	
0x0007	
0x0008	
0x0009	
...	

- memory is organized sequentially
- memory is addressed starting from the *top*
 - granularity is a byte

- a process is an instance of a computer program
- a multitasking operating system is capable of running many processes despite having only one or a “small” number of physical cores
 - context switches the processes at predefined time slices
 - **top** activity

memory organization



- memory is organized sequentially
- memory is addressed starting from the *top*
 - granularity is a byte
- a process and its data can be stored anywhere in memory

- a process is an instance of a computer program
- a multitasking operating system is capable of running many processes despite having only one or a “small” number of physical cores
 - context switches the processes at predefined time slices
 - **top** activity

memory organization

0x0000	
0x0001	
0x0002	
0x0003	
0x0004	
0x0005	
0x0006	
0x0007	
0x0008	
0x0009	
...	

- memory is organized sequentially
- memory is addressed starting from the *top*
 - granularity is a byte
- a process and its data can be stored anywhere in memory
 - physical vs. logical addressing

- a process is an instance of a computer program
- a multitasking operating system is capable of running many processes despite having only one or a “small” number of physical cores
 - context switches the processes at predefined time slices
 - **top** activity

memory organization

0x0000	
0x0001	
0x0002	
0x0003	
0x0004	
0x0005	
0x0006	
0x0007	
0x0008	
0x0009	
...	

- memory is organized sequentially
- memory is addressed starting from the *top*
 - granularity is a byte
- a process and its data can be stored anywhere in memory
 - physical vs. logical addressing
 - each process has its own *address space*

- a process is an instance of a computer program
- a multitasking operating system is capable of running many processes despite having only one or a “small” number of physical cores
 - context switches the processes at predefined time slices
 - **top** activity

- what is it?

- `lscpu`
- `cat /proc/cpuinfo` activity for CPU info
- `cat /sys/devices/system/cpu/cpu0/cache/` for cache info
- L1, L2, L3
 - L1 is usually split into data and instruction
 - L2 is usually not shared, i.e., each core has its own
 - L3 is usually shared by multiple cores

cache memory

- a high-speed memory co-located with the CPU on the chip

- `lscpu`
- `cat /proc/cpuinfo` activity for CPU info
- `cat /sys/devices/system/cpu/cpu0/cache/` for cache info
- L1, L2, L3
 - L1 is usually split into data and instruction
 - L2 is usually not shared, i.e., each core has its own
 - L3 is usually shared by multiple cores

cache memory

- a high-speed memory co-located with the CPU on the chip
- what problem is cache addressing?

- `lscpu`
- `cat /proc/cpuinfo` activity for CPU info
- `cat /sys/devices/system/cpu/cpu0/cache/` for cache info
- L1, L2, L3
 - L1 is usually split into data and instruction
 - L2 is usually not shared, i.e., each core has its own
 - L3 is usually shared by multiple cores

cache memory

- a high-speed memory co-located with the CPU on the chip
- CPU can compute faster than memory can be fetched
 - many problems reuse data or use data that is stored near each other in memory, known as *temporal* and *spatial* locality, respectively

- `lscpu`
- `cat /proc/cpuinfo` activity for CPU info
- `cat /sys/devices/system/cpu/cpu0/cache/` for cache info
- L1, L2, L3
 - L1 is usually split into data and instruction
 - L2 is usually not shared, i.e., each core has its own
 - L3 is usually shared by multiple cores

cache memory

- a high-speed memory co-located with the CPU on the chip
- CPU can compute faster than memory can be fetched
 - many problems reuse data or use data that is stored near each other in memory, known as *temporal* and *spatial* locality, respectively
- how does it work?

- `lscpu`
- `cat /proc/cpuinfo` activity for CPU info
- `cat /sys/devices/system/cpu/cpu0/cache/` for cache info
- L1, L2, L3
 - L1 is usually split into data and instruction
 - L2 is usually not shared, i.e., each core has its own
 - L3 is usually shared by multiple cores

cache memory

- a high-speed memory co-located with the CPU on the chip
- CPU can compute faster than memory can be fetched
 - many problems reuse data or use data that is stored near each other in memory, known as *temporal* and *spatial* locality, respectively
- when memory load is issued, cache is checked to see if the data exists there. if not, it is loaded from main memory and stored in cache. when cache becomes full, old data is evicted to make room for new.

- `lscpu`
- `cat /proc/cpuinfo` activity for CPU info
- `cat /sys/devices/system/cpu/cpu0/cache/` for cache info
- L1, L2, L3
 - L1 is usually split into data and instruction
 - L2 is usually not shared, i.e., each core has its own
 - L3 is usually shared by multiple cores

cache memory

- a high-speed memory co-located with the CPU on the chip
- CPU can compute faster than memory can be fetched
 - many problems reuse data or use data that is stored near each other in memory, known as *temporal* and *spatial* locality, respectively
- when memory load is issued, cache is checked to see if the data exists there. if not, it is loaded from main memory and stored in cache. when cache becomes full, old data is evicted to make room for new.
- how much is there?

- `lscpu`
- `cat /proc/cpuinfo` activity for CPU info
- `cat /sys/devices/system/cpu/cpu0/cache/` for cache info
- L1, L2, L3
 - L1 is usually split into data and instruction
 - L2 is usually not shared, i.e., each core has its own
 - L3 is usually shared by multiple cores

cache memory

- a high-speed memory co-located with the CPU on the chip
- CPU can compute faster than memory can be fetched
 - many problems reuse data or use data that is stored near each other in memory, known as *temporal* and *spatial* locality, respectively
- when memory load is issued, cache is checked to see if the data exists there. if not, it is loaded from main memory and stored in cache. when cache becomes full, old data is evicted to make room for new.
- multiple levels of cache, each larger than the previous, but all much smaller than main memory


- `lscpu`
- `cat /proc/cpuinfo` activity for CPU info
- `cat /sys/devices/system/cpu/cpu0/cache/` for cache info
- L1, L2, L3
 - L1 is usually split into data and instruction
 - L2 is usually not shared, i.e., each core has its own
 - L3 is usually shared by multiple cores

From Intel Performance Analysis Guide:

Core i7 Xeon 5500 Series Data Source Latency (approximate) [Pg. 22]

local	L1 CACHE hit,	~4 cycles (2.1 - 1.2 ns)
local	L2 CACHE hit,	~10 cycles (5.3 - 3.0 ns)
local	L3 CACHE hit, line unshared	~40 cycles (21.4 - 12.0 ns)
local	L3 CACHE hit, shared line in another core	~65 cycles (34.8 - 19.5 ns)
local	L3 CACHE hit, modified in another core	~75 cycles (40.2 - 22.5 ns)
remote	L3 CACHE (Ref: Fig.1 [Pg. 5])	~100-300 cycles (160.7 - 30.0 ns)
local	DRAM	~60 ns
remote	DRAM	~100 ns

more about latency



Latency Numbers Every Programmer Should Know



except where otherwise noted, this worked is licensed under creative commons attribution-sharealike 4.0 international license