


Introduction to MPI

Jeremy Iverson

College of Saint Benedict & Saint John's University



- A standard for explicit distributed memory parallel computation.
- Many implementations available, both open-source and proprietary.

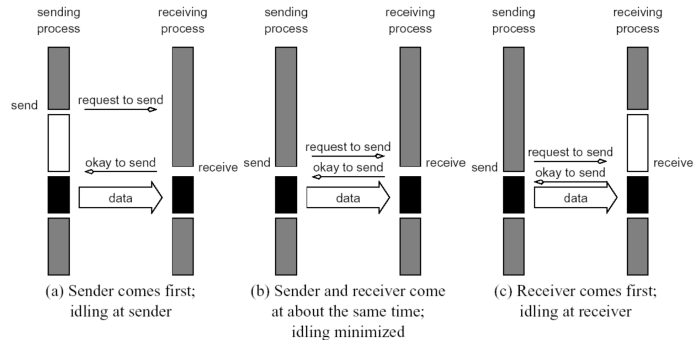
1. We are using an implementation of MPI called Open-MPI

- Uses the *SPMD* model of parallelism
 - All processes execute the same program.
 - Different processes carry out different actions by conditional execution of code based on processes' rank.
 - Processes can communicate with each other by sending explicit messages

- Requires inclusion of `mpi.h` header file.
- `MPI_Init()`
- `MPI_Finalize()`
- `MPI_Comm_size()`
- `MPI_Comm_rank()`
- `MPI_Send()`
- `MPI_Recv()`

point-to-point communication

- MPI uses *communicators* to organize processes. Processes can only communicate with other processes in the same communicator. The base communicator to which all processes belong is called **MPI_COMM_WORLD**.
- Programs can deadlock due to improperly ordered or unmatched point-to-point communications.



collective communication

- Represent regular *communication patterns* that are performed by parallel algorithms.
- Include groups of processes, not just two.
- Can be implemented to take advantage of underlying network characteristics and thus improve performance compared to simple point-to-point equivalents.
- Most parallel libraries provide functions to perform them (`omp parallel for reduction(+:sum)`)

```
1 if (0 == rank) {  
2     for (int r = 1; r < p; r++) {  
3         MPI_Send(&x, 1, MPI_INT, r, 0,  
4             MPI_COMM_WORLD);  
5     }  
6 } else {  
7     MPI_Recv(&x, 1, MPI_INT, 0, 0,  
8         MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
9 }
```

```
1 if (0 == rank) {  
2     for (int r = 1; r < p; r++) {  
3         MPI_Send(&x, 1, MPI_INT, r, 0,  
4             MPI_COMM_WORLD);  
5     }  
6 } else {  
7     MPI_Recv(&x, 1, MPI_INT, 0, 0,  
8         MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
9 }
```

```
1 MPI_Bcast(&x, 1, MPI_INT, 0, MPI_COMM_WORLD);
```



```
1 if (0 == rank) {  
2     for (int r = 1; r < p; r++) {  
3         MPI_Send(rating + r * n * m, n * m,  
4                 MPI_DOUBLE, r, 0, MPI_COMM_WORLD);  
5     }  
6 } else {  
7     MPI_Recv(rating, n * m, MPI_DOUBLE, 0, 0,  
8             MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
9 }
```

```
1 if (0 == rank) {  
2     for (int r = 1; r < p; r++) {  
3         MPI_Send(rating + r * n * m, n * m,  
4                 MPI_DOUBLE, r, 0, MPI_COMM_WORLD);  
5     }  
6 } else {  
7     MPI_Recv(rating, n * m, MPI_DOUBLE, 0, 0,  
8             MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
9 }
```

```
1 MPI_Scatter(rating, n * m, MPI_DOUBLE,  
2 rating, n * m, MPI_DOUBLE, 0,  
3 MPI_COMM_WORLD);
```

```
1 if (0 == rank) {  
2     for (int r = 1; r < p; r++) {  
3         size_t const rn = (r + 1) * base > n  
4                             ? n - r * base : base;  
5         MPI_Send(rating + r * base * m, rn * m,  
6                 MPI_DOUBLE, r, 0, MPI_COMM_WORLD);  
7     }  
8 } else {  
9     MPI_Recv(rating, ln * m, MPI_DOUBLE, 0, 0,  
10            MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
11 }
```

```
1 if (0 == rank) {  
2     for (int r = 1; r < p; r++) {  
3         size_t const rn = (r + 1) * base > n  
4                             ? n - r * base : base;  
5         MPI_Send(rating + r * base * m, rn * m,  
6                 MPI_DOUBLE, r, 0, MPI_COMM_WORLD);  
7     }  
8 } else {  
9     MPI_Recv(rating, ln * m, MPI_DOUBLE, 0, 0,  
10            MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
11 }
```

```
1 MPI_Scatterv(...);
```



except where otherwise noted, this worked is licensed under creative commons attribution-sharealike 4.0 international license