

Performance analysis

Jeremy Iverson

College of Saint Benedict & Saint John's University

Performance analysis

- how do we reason about parallel algorithms?
- how can we compare two algorithms and determine which is better?
- how do we measure improvement?

Performance metrics

- execution time (T_p)
- speedup (S)
- efficiency (E)
- cost (C)

Execution time

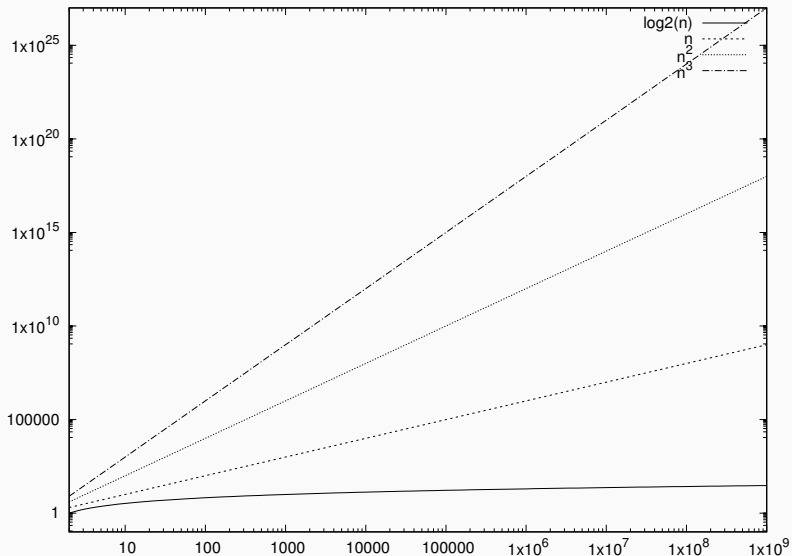
Serial (T_s)

- time elapsed between beginning and end of execution

Parallel (T_p)

- time elapsed between beginning of execution and the moment the last processing element finishes execution
- Adding numbers
- Dot-product
- Matrix-vector multiplication
- Matrix-matrix multiplication

Execution time



Speedup


Speedup ($S = T_s/T_p$)

- the ratio of time taken to solve a problem on a single processing element to the time required to solve the same problem on a parallel computer with p processing elements

Speedup

$$\text{Speedup } (S = T_s/T_p)$$

- the ratio of time taken to solve a problem on a single processing element to the time required to solve the same problem on a parallel computer with p processing elements



Efficiency


Efficiency ($E = S/p$)

- the ratio of speedup to the number of processing elements — the fraction of time for which a processing element is usefully employed

Efficiency

Efficiency ($E = S/p$)

- the ratio of speedup to the number of processing elements — the fraction of time for which a processing element is usefully employed




Cost ($C = pT_p$)

- the sum of the time spent by all processing elements solving the problem
- *cost-optimal* if $C = T_s$

Cost ($C = pT_p$)

- the sum of the time spent by all processing elements solving the problem
- *cost-optimal* if $C = T_s$



Exercise — vector summation

$p = n$ — *not cost-optimal*

- $T_p = \Theta(\log n)$
- $S = \Theta(\frac{n}{\log n})$
- $E = \Theta(\frac{1}{\log n})$
- $C = \Theta(n \log n)$

$p > n$ — too many processing elements, use less

$p < n$ — ?

Exercise — vector summation

$p = n$ — not cost-optimal

- $T_p = \Theta(\log n)$
- $S = \Theta(\frac{n}{\log n})$
- $E = \Theta(\frac{1}{\log n})$
- $C = \Theta(n \log n)$

$p > n$ — too many processing elements, use less

$p < n$ — not cost-optimal?

- $T_p = \Theta(\frac{n}{p} + \log p)$
- $S = \Theta(\frac{n}{\frac{n}{p} + \log p})$
- $E = \Theta(\frac{n}{n + \log p})$
- $C = \Theta(n + p \log p)$

Exercise — vector summation

$p = n$ — not cost-optimal


- $T_p = \Theta(\log n)$
- $S = \Theta(\frac{n}{\log n})$
- $E = \Theta(\frac{1}{\log n})$
- $C = \Theta(n \log n)$

$p > n$ — too many processing elements, use less

$p < n$ — cost-optimal iff $n = \Theta(p \log p)$

- $T_p = \Theta(\frac{n}{p} + \log p)$
- $S = \Theta(\frac{n}{\frac{n}{p} + \log p})$
- $E = \Theta(\frac{n}{n + \log p})$
- $C = \Theta(n + p \log p)$

Exercise



Counting sort

```
1: function COUNTINGSORT(inValues,outValues)
2:   count  $\leftarrow \{0\}$ 
3:   for all v in inValues do
4:     count[v]  $\leftarrow$  count[v] + 1
5:   end for
6:   PREFIXSCAN(count)
7:   for all v in inValues do
8:     outValues[count[v]]  $\leftarrow$  v
9:     count[v]  $\leftarrow$  count[v] + 1
10:  end for
11: end function
```


Counting sort — analysis

$$p = n - ?$$

- $T_p = \Theta(\log |count|)$ or $T_p = \Theta(\log n)$
- $S = \Theta(\frac{n}{\log n})$
- $E = \Theta(\frac{1}{\log n})$
- $C = \Theta(n \log n)$



except where otherwise noted, this worked is licensed under creative commons attribution-sharealike 4.0 international license