

# Performance analysis

---

Jeremy Iverson

College of Saint Benedict & Saint John's University

## Performance analysis

- how do we reason about parallel algorithms?
- how can we compare two algorithms and determine which is better?
- how do we measure improvement?

## Performance metrics

- execution time ( $T_p$ )
- speedup ( $S$ )
- efficiency ( $E$ )
- cost ( $C$ )

# Execution time

## Serial ( $T_s$ )


- time elapsed between beginning and end of execution

## Parallel ( $T_p$ )

- time elapsed between beginning of execution and the moment the last processing element finishes execution
- Adding numbers
- Dot-product
- Matrix-vector multiplication
- Matrix-matrix multiplication

1. Adding numbers —  $T_s = \Theta(n)$  —  $T_p = \Theta(\log n)$
2. Matrix-vector —  $T_s = \Theta(n^2)$  —  $T_p = \Theta(n)$
3. Matrix-matrix —  $T_s = \Theta(n^3)$  —  $T_p = \Theta(n^2)$
4. So how long to compute mat-mat for a  $10000 \times 10000$  mat?

# Execution time



# Speedup

**Speedup** ( $S = T_s/T_p$ )


- the ratio of time taken to solve a problem on a single processing element to the time required to solve the same problem on a parallel computer with  $p$  processing elements

1. Adding numbers —  $S = \Theta(n/\log n)$
2. Matrix-vector —  $S = \Theta(n^2/n) = \Theta(n)$
3. Matrix-matrix —  $S = \Theta(n^3/n^2) = \Theta(n)$
4. What are the limits of speedup?

# Speedup

Speedup ( $S = T_s/T_p$ )

- the ratio of time taken to solve a problem on a single processing element to the time required to solve the same problem on a parallel computer with  $p$  processing elements



1. Adding numbers —  $S = \Theta(n/\log n)$
2. Matrix-vector —  $S = \Theta(n^2/n) = \Theta(n)$
3. Matrix-matrix —  $S = \Theta(n^3/n^2) = \Theta(n)$
4. What are the limits of speedup?

# Efficiency

## Efficiency ( $E = S/p$ )

- the ratio of speedup to the number of processing elements — the fraction of time for which a processing element is usefully employed


1. Adding numbers —  $E = \Theta(n/\log n)/n = \Theta(1/\log n)$
2. Matrix-vector —  $E = \Theta(n)/n = \Theta(1)$
3. Matrix-matrix —  $E = \Theta(n)/n = \Theta(1)$
4. Why do you think that the efficiency of adding numbers is less than matrix-vector or matrix-matrix?
5. What are the limits of efficiency?



# Efficiency

Efficiency ( $E = S/p$ )

- the ratio of speedup to the number of processing elements — the fraction of time for which a processing element is usefully employed



- Adding numbers —  $E = \Theta(n/\log n)/n = \Theta(1/\log n)$
- Matrix-vector —  $E = \Theta(n)/n = \Theta(1)$
- Matrix-matrix —  $E = \Theta(n)/n = \Theta(1)$
- Why do you think that the efficiency of adding numbers is less than matrix-vector or matrix-matrix?
- What are the limits of efficiency?


**Cost** ( $C = pT_p$ )

- the sum of the time spent by all processing elements solving the problem
- *cost-optimal* if  $C = T_s$

1. Adding numbers —  $C = \Theta(n \log n)$
2. Matrix-vector —  $C = n \times \Theta(n) = \Theta(n^2)$
3. Matrix-matrix —  $C = n \times \Theta(n^2) = \Theta(n^3)$
4. Adding numbers is not cost-optimal, the others are

Cost ( $C = pT_p$ )

- the sum of the time spent by all processing elements solving the problem
- *cost-optimal* if  $C = T_s$



1. Adding numbers —  $C = \Theta(n \log n)$
2. Matrix-vector —  $C = n \times \Theta(n) = \Theta(n^2)$
3. Matrix-matrix —  $C = n \times \Theta(n^2) = \Theta(n^3)$
4. Adding numbers is not cost-optimal, the others are

## Exercise — vector summation

$p = n$  — *not cost-optimal*

- $T_p = \Theta(\log n)$
- $S = \Theta(\frac{n}{\log n})$
- $E = \Theta(\frac{1}{\log n})$
- $C = \Theta(n \log n)$

$p > n$  — too many processing elements, use less

$p < n$  — ?

## Exercise — vector summation

$p = n$  — not cost-optimal

- $T_p = \Theta(\log n)$
- $S = \Theta(\frac{n}{\log n})$
- $E = \Theta(\frac{1}{\log n})$
- $C = \Theta(n \log n)$

$p > n$  — too many processing elements, use less

$p < n$  — not cost-optimal?

- $T_p = \Theta(\frac{n}{p} + \log p)$
- $S = \Theta(\frac{n}{\frac{n}{p} + \log p})$
- $E = \Theta(\frac{n}{n + \log p})$
- $C = \Theta(n + p \log p)$

## Exercise — vector summation

$p = n$  — not cost-optimal

- $T_p = \Theta(\log n)$
- $S = \Theta(\frac{n}{\log n})$
- $E = \Theta(\frac{1}{\log n})$
- $C = \Theta(n \log n)$


$p > n$  — too many processing elements, use less

$p < n$  — cost-optimal iff  $n = \Theta(p \log p)$

- $T_p = \Theta(\frac{n}{p} + \log p)$
- $S = \Theta(\frac{n}{\frac{n}{p} + \log p})$
- $E = \Theta(\frac{n}{n + \log p})$
- $C = \Theta(n + p \log p)$

1. if you are given a problem size, what is the maximum number of processing elements that can be used in a cost optimal way?
2. what are the limits of efficiency...how about speedup?

## Exercise



1. give students introduction to Gnuplot

## Counting sort

```
1: function COUNTINGSORT(inValues,outValues)
2:   count  $\leftarrow \{0\}$ 
3:   for all v in inValues do
4:     count[v]  $\leftarrow$  count[v] + 1
5:   end for
6:   PREFIXSCAN(count)
7:   for all v in inValues do
8:     outValues[count[v]]  $\leftarrow$  v
9:     count[v]  $\leftarrow$  count[v] + 1
10:  end for
11: end function
```

1. Which of these steps can be parallelized?
2. Assuming we have  $p = n$ , and we can efficiently parallelize the two for loops, what is our analysis?



## Counting sort — analysis

$$p = n - ?$$

- $T_p = \Theta(\log |count|)$  or  $T_p = \Theta(\log n)$
- $S = \Theta(\frac{n}{\log n})$
- $E = \Theta(\frac{1}{\log n})$
- $C = \Theta(n \log n)$

1.  $T_p$  will be  $\log |count|$  unless we use a dense data structure like a map in which case it will be  $\log n$
2.  $n$  will always be  $\leq |count|$



except where otherwise noted, this worked is licensed under creative commons attribution-sharealike 4.0 international license